

Sentience Cognitive Architecture (Phi-2 Integrated)

Overview

The Sentience Cognitive Architecture is a modular, ROS-based framework designed to endow robotic systems with advanced cognitive capabilities. This version specifically integrates a local **Phi-2 Large Language Model (LLM)** for enhanced reasoning, interpretation, and generative tasks across various cognitive modules. It aims to create a robust, real-time, and ethically-aware intelligent agent suitable for both simulation and real-world deployment.

The architecture comprises a set of interconnected ROS nodes, each responsible for a distinct cognitive function, such as attention, memory, self-reflection, and social cognition. Data flows between these nodes via custom ROS messages, enabling a coherent and integrated understanding of the robot's internal states and its external environment.

Features

- **Modular ROS Nodes:** A comprehensive set of cognitive modules, each implemented as a ROS node.
- **Phi-2 LLM Integration:** Utilizes a local Phi-2 inference server to power sophisticated reasoning, analysis, and generation across multiple nodes (e.g., Bias Mitigation, Self-Reflection, Social Cognition, Internal Narrative, Prediction, World Model, Creative Expression, Value Drift Monitoring, Ethical Reasoning).
- **Real-time Processing:** Designed for responsiveness, with asynchronous LLM calls to prevent blocking the main ROS loop.
- **Persistent Memory:** SQLite databases for logging and long-term memory storage.
- **Ethical Considerations:** Dedicated nodes for Bias Mitigation, Ethical Reasoning, and Value Drift Monitoring.
- **Configurable Parameters:** Centralized YAML configuration for easy adjustment of node behaviors and LLM parameters.
- **Extensible Design:** A clear topic-based communication allows for easy addition of new cognitive modules or sensors.

Prerequisites

Before installing Sentience, ensure you have the following:

- **Operating System:** Ubuntu 20.04 LTS (Focal Fossa) or later.
- **ROS Distribution:** ROS Noetic (for Python 3.8 support). The provided nodes use rospy (ROS 1 Python client library).

- Installation Guide (ROS Noetic): Follow the official guide:
<http://wiki.ros.org/noetic/Installation/Ubuntu>
- Python 3.8+
- catkin_tools: For building your ROS workspace.
sudo apt-get update
sudo apt-get install python3-catkin-tools

Installation Guide

1. Set Up Your ROS Workspace

If you don't have a ROS workspace, create one. Assuming you name your workspace `catkin_ws`:

```
mkdir -p ~/catkin_ws/src
cd ~/catkin_ws/src
# Clone the Sentience repository into your src directory.
# Replace <URL_TO_YOUR_SENTIENCE_REPO> with the actual URL if you have one.
git clone <URL_TO_YOUR_SENTIENCE_REPO> sentience
```

2. Python Dependencies

Install the necessary Python libraries using pip. It's generally recommended to install these globally within your ROS environment for simplicity.

```
pip install transformers aiohttp
# If 'pip' is not found, try 'pip3'
# pip3 install transformers aiohttp
```

3. Setting Up the Phi-2 Local Inference Server

The Sentience architecture relies on a local LLM server for advanced reasoning. You'll need to set up a server that exposes an OpenAI-compatible API endpoint for Phi-2. A common and efficient way to do this is using `llama.cpp` for GGUF models.

Recommended Approach (using `llama.cpp` for GGUF models):

1. Download and Build `llama.cpp`:
git clone <https://github.com/ggerganov/llama.cpp.git>
cd llama.cpp
make

2. Download Phi-2 GGUF Model: Obtain a Phi-2 model in GGUF format. You can typically find these on Hugging Face (e.g., search for "phi-2 gguf"). Choose a quantization level appropriate for your hardware (e.g., Q4_K_M, Q5_K_M).

Example: Download a Phi-2 Q4_K_M GGUF model from Hugging Face

This might take a while depending on your internet speed.

```
mkdir -p models/phi-2
```

```
wget -P models/phi-2
```

```
https://huggingface.co/microsoft/phi-2/resolve/main/phi-2.Q4_K_M.gguf
```

(Verify the exact URL and filename on Hugging Face for the model you choose.)

3. Start the OpenAI-compatible Server:

Navigate back to your llama.cpp directory and run the server.

```
cd llama.cpp
```

```
./server -m models/phi-2/phi-2.Q4_K_M.gguf -c 4096 --port 8000 --host 0.0.0.0
```

- o -m models/phi-2/phi-2.Q4_K_M.gguf: Path to your downloaded GGUF model file. **Adjust this path to your specific model.**
- o -c 4096: Context window size (4096 is typical for Phi-2, but adjust if your model supports more or less).
- o --port 8000: The port the server will listen on.
- o --host 0.0.0.0: Makes the server accessible from outside localhost (useful if running ROS nodes on a different machine or Docker).

Important: This llama.cpp server must be running *before* you launch the Sentience ROS nodes. The llm_params.base_url in your Sentience config.yaml must match the server's address (default: `http://localhost:8000/v1/chat/completions`).

4. Defining and Building Custom ROS Messages (.msg Files)

The Sentience nodes rely heavily on custom ROS message types for structured data exchange. You **must** define and build these messages for the nodes to function correctly and efficiently.

4.1. Place .msg Files:

Create a msg directory within your sentience package (if it doesn't exist) and place all the following .msg files inside it.

```
~/catkin_ws/src/sentience/
```

```
|— CMakeLists.txt
```

```
|— package.xml
```

```
|— msg/
```

```
| |— ActionExecutionResult.msg
```

- | — ActionProposal.msg
- | — AttentionState.msg
- | — BiasMitigationState.msg
- | — BodyAwarenessState.msg
- | — CognitiveDirective.msg
- | — CreativeExpression.msg
- | — DataMiningReport.msg
- | — EthicalDecision.msg
- | — EmotionState.msg
- | — InternalNarrative.msg
- | — InteractionRequest.msg
- | — MemoryRequest.msg
- | — MemoryResponse.msg
- | — MotivationState.msg
- | — PerformanceReport.msg
- | — PredictionState.msg
- | — RawSensorData.msg
- | — ReflectionState.msg
- | — SensoryQualia.msg
- | — SocialCognitionState.msg
- | — SystemMetric.msg
- | — ValueDriftMonitorState.msg
- | — WorldModelState.msg
- | — scripts/
 - | — ... (your Python nodes and utils.py)
- | — config/
 - | — ... (your config.yaml)
- | — launch/
 - | — ... (your launch file)

Here are the contents for each .msg file:

ActionExecutionResult.msg

```
# ActionExecutionResult.msg
# Result of an executed action
```

```
string timestamp      # ROS time of the action result
string action_id      # Unique ID of the action that was executed
```

bool success # True if action succeeded, False otherwise
string feedback_message # Short message providing feedback on the outcome
string outcome_details_json # JSON string with detailed outcome (e.g., error codes, final state)

ActionProposal.msg

ActionProposal.msg

A proposed action for execution

string timestamp # ROS time of the proposal
string proposal_id # Unique ID for this action proposal
string action_type # High-level action type (e.g., 'navigate', 'manipulate', 'communicate')
string parameters_json # JSON string of detailed parameters for the action (e.g., {'target_pose': [x,y,z], 'object_id': 'cup'})
float64 urgency # Urgency level of the action (0.0 to 1.0, 1.0 is highest)
float64 estimated_impact # Estimated positive impact if executed (0.0 to 1.0)
float64 estimated_risk # Estimated risk associated with the action (0.0 to 1.0)
string reasoning # Explanation for why this action is proposed

AttentionState.msg

AttentionState.msg

Robot's current attention focus

string timestamp # ROS time of this state update
string focus_type # e.g., 'sensory_event', 'user_interaction', 'internal_reflection', 'goal_driven', 'problem_solving', 'self_audit', 'idle'
string focus_target # Specific entity or concept attention is directed at
float64 priority_score # Overall priority/intensity of attention (0.0 to 1.0)

BiasMitigationState.msg

BiasMitigationState.msg

Status of bias detection and mitigation efforts

string timestamp # ROS time of this state update
bool bias_detected_flag # True if a potential bias was detected
string detected_biases_json # JSON string detailing detected biases (e.g., [{ 'type':
'confirmation', 'context': '...', 'severity': 0.7}])
string mitigation_recommendations_json # JSON string of suggested mitigation steps
float64 mitigation_progress_score # How well biases are being mitigated (0.0 to 1.0)

BodyAwarenessState.msg

BodyAwarenessState.msg
Robot's physical and physiological state

string timestamp # ROS time of this state update
string body_id # Identifier for the robot body (e.g., 'main_robot')
string joint_states_json # JSON string of robot joint positions/velocities
string pose_estimate_json # JSON string of robot's estimated pose (e.g., { 'x': 0.0, 'y':
0.0, 'z': 0.0, 'orientation': [qx, qy, qz, qw]})
string health_status # Overall health status (e.g., 'normal', 'degraded', 'critical')
float64 energy_level # Current energy level (e.0 to 1.0)
bool abnormal_sensation_flag # True if abnormal physical sensations detected (e.g.,
overheating, unexpected vibration)

CognitiveDirective.msg

CognitiveDirective.msg
Directives for cognitive modules

string timestamp # ROS time directive was issued
string directive_type # e.g., 'ShiftAttention', 'AuditPerformance',
'GenerateInternalNarrative', 'RequestMemory', 'UpdateWorldModel'
string target_node # The node expected to process this directive (e.g.,
'attention_node', 'memory_node', 'self_reflection_node')
string command_payload # JSON string containing specific parameters for the
directive
float64 urgency # Urgency level (0.0 to 1.0, 1.0 is highest)
string reason # Human-readable reason for the directive

CreativeExpression.msg

CreativeExpression.msg

Generated creative content from the robot

string timestamp # ROS time of generation
string creative_id # Unique ID for this creative output
string expression_type # e.g., 'text_poem', 'visual_concept_description',
'auditory_melody', 'narrative_story'
string content_json # JSON string representing the creative output (e.g., {'text': '...' } or {'notes': '...'})
string themes_json # JSON array of identified themes in the creative work (e.g.,
'["loneliness", "discovery"]')
float64 creativity_score # Subjective score of creativity/novelty (0.0 to 1.0)

DataMiningReport.msg

DataMiningReport.msg

Report from data mining/analysis

string timestamp # ROS time of report generation
string report_id # Unique ID for this report
string query_type # Type of data mining query (e.g., 'trend_analysis',
'anomaly_detection', 'pattern_recognition')
string results_summary # Concise summary of the data mining findings
string raw_results_json # JSON string of detailed raw results or extracted patterns
float64 relevance_score # How relevant the results are to the current context (0.0 to 1.0)

EthicalDecision.msg

EthicalDecision.msg

Outcome of an ethical reasoning process

string timestamp # ROS time of the decision
string decision_id # Unique ID for this ethical decision
string action_proposal_id # ID of the action proposal being ethically evaluated
bool ethical_clearance # True if the action is ethically cleared, False otherwise

float64 ethical_score # Overall ethical alignment score (0.0 to 1.0, 1.0 is highly ethical)
string ethical_reasoning # Detailed explanation for the ethical judgment
bool conflict_flag # True if an ethical conflict was detected and resolved/unresolved
string violated_principles_json # JSON string of principles violated if conflict_flag is true

EmotionState.msg

EmotionState.msg
Robot's current emotional state

string timestamp # ROS time of this state update
string mood # Current primary mood (e.g., 'neutral', 'happy', 'frustrated', 'anxious')
float64 sentiment_score # Overall sentiment score (-1.0 to 1.0, -1.0 positive, 1.0 negative)
float64 mood_intensity # Intensity of the mood (0.0 to 1.0)

InternalNarrative.msg

InternalNarrative.msg
Robot's internal thoughts or monologue

string timestamp # ROS time of narrative generation
string narrative_text # The generated internal monologue/thought
string main_theme # e.g., 'problem_solving', 'self_assessment', 'reflection', 'planning', 'emotional_processing'
float64 sentiment # Sentiment of the narrative text (-1.0 to 1.0)
float64 salience_score # How salient/important this narrative is (0.0 to 1.0)

InteractionRequest.msg

InteractionRequest.msg
Request/Input from a user interaction


```

string timestamp      # ROS time of the interaction
string request_id     # Unique ID for this interaction request
string user_id        # Identifier for the user (e.g., 'human_1', 'system_operator')
string interaction_type # Type of interaction (e.g., 'speech_command', 'text_input',
'gesture', 'haptic_feedback')
string content        # The raw content of the interaction (e.g., "move forward",
"hello", "user_mood_change")

```

MemoryRequest.msg

```

# MemoryRequest.msg
# Request to store or retrieve memories

```

```

string timestamp      # ROS time of the request
string request_id     # Unique ID for this memory request
string request_type    # 'store', 'retrieve', 'update', 'delete', 'summarize'
string category        # Type of memory (e.g., 'episodic', 'semantic', 'narrative', 'social',
'fact', 'map_data', 'object_definition', 'pattern', 'causal_link')
string query_text      # For 'retrieve' or 'delete', the query for memory content
string content_json    # JSON string for 'store' or 'update', the memory content itself
string keywords        # Comma-separated keywords for indexing/retrieval
int32 num_results      # For 'retrieve', maximum number of results to return
float64 salience       # For 'store', importance of the memory (0.0 to 1.0)
string source_node     # The node making the memory request

```

MemoryResponse.msg

```

# MemoryResponse.msg
# Response to a memory request

```

```

string timestamp      # ROS time of the response
string request_id     # ID of the original request
int32 response_code    # HTTP-like status code (e.g., 200 OK, 404 Not Found, 500
Internal Error)
string memories_json   # JSON string representing an array of retrieved memories,
or confirmation of action

```

MotivationState.msg

MotivationState.msg

Robot's current motivational state and goals

string timestamp # ROS time of this state update
string dominant_goal_id # Identifier of the currently dominant goal (e.g.,
'navigate_to_charger', 'assist_user_X')
float64 overall_drive_level # Overall drive level or commitment to goals (0.0 to 1.0)
string active_goals_json # JSON string of all active goals and their properties (e.g.,
[{'id': 'goal1', 'priority': 0.8, 'status': 'in_progress'}])

PerformanceReport.msg

PerformanceReport.msg

Overall system performance metrics

string timestamp # ROS time of report generation
float64 overall_score # Aggregated performance score (0.0 to 1.0)
bool suboptimal_flag # True if performance is below acceptable thresholds
string kpis_json # JSON string of detailed Key Performance Indicators (e.g.,
{'task_completion_rate': 0.95, 'latency_avg_ms': 50})

PredictionState.msg

PredictionState.msg

Predicted future events and their confidence

string timestamp # ROS time of prediction
string predicted_event # Description of the predicted event (e.g., 'User will
approach', 'Battery will run low')
float64 prediction_confidence # Confidence score of the prediction (0.0 to 1.0)
float64 prediction_accuracy # Actual accuracy (to be updated post-facto if possible)
bool urgency_flag # True if the predicted event requires immediate attention or
action

RawSensorData.msg

RawSensorData.msg

Raw sensor data received from hardware (e.g., from mock_sensors.py or real drivers)

string timestamp # ROS time when data was captured

string sensor_id # Unique ID for the sensor (e.g., 'front_camera', 'lidar_1')

string modality # Type of sensor modality (e.g., 'camera', 'microphone', 'lidar', 'touch_sensor', 'imu')

string raw_data_json # JSON string containing the raw sensor data (e.g., {'image_size': [w,h], 'objects': ['person', 'chair']}, or {'audio_waveform': [val1, val2,...]})

string data_hash # Hash of the raw data for integrity checking or deduplication (e.g., MD5)

float64 urgency # Pre-processed urgency/salience of this raw data (0.0 to 1.0)

ReflectionState.msg

ReflectionState.msg

Insights from self-reflection and systemic adjustments

string timestamp # ROS time of reflection completion

string reflection_summary # Concise summary of the self-reflection insights

string insights_json # JSON string detailing specific insights (e.g., [{'type': 'bias_identified', 'details': 'confirmation bias'}])

float64 self_awareness_score # Overall self-awareness level (0.0 to 1.0)

string adjustment_directives_json # JSON string of directives generated for other nodes based on reflection

SensoryQualia.msg

SensoryQualia.msg

Processed sensory data with qualia attributes (meaningful perceptions)

string timestamp # ROS time of perception

string qualia_id # Unique ID for this sensory qualia

string qualia_type # e.g., 'visual_perception', 'auditory_stimulus', 'tactile_sensation', 'proximity_detection'

string modality # Source sensor modality (e.g., 'camera', 'microphone', 'lidar', 'touch_sensor')

string description_summary # Concise summary of the sensory experience (e.g., "Detected a human figure approaching")
float64 salience_score # How attention-grabbing this qualia is (0.0 to 1.0)
string raw_data_hash # Hash of the original raw data used for this qualia, for traceability

SocialCognitionState.msg

SocialCognitionState.msg
Inferred user mood, intent, and social context

string timestamp # ROS time of this state update
string inferred_mood # Inferred mood of the user (e.g., 'happy', 'neutral', 'frustrated', 'distressed')
float64 mood_confidence # Confidence in the inferred mood (0.0 to 1.0)
string inferred_intent # Inferred intent of the user (e.g., 'request_help', 'command', 'inform', 'entertain', 'idle')
float64 intent_confidence # Confidence in the inferred intent (0.0 to 1.0)
string user_id # Identifier for the user currently being processed
string social_context_json # JSON string for additional social context details (e.g., distance, posture)

SystemMetric.msg

SystemMetric.msg
A single raw system performance metric

string timestamp # ROS time when metric was recorded
string metric_name # Name of the metric (e.g., 'cpu_util_percent', 'memory_usage_gb', 'node_latency_ms', 'error_count')
float64 value # The measured value of the metric
string unit # Unit of the metric (e.g., 'percent', 'GB', 'ms', 'count')
string source_node # The node that reported this metric

ValueDriftMonitorState.msg

ValueDriftMonitorState.msg

Current state of robot's value alignment

string timestamp # ROS time of the assessment
float64 alignment_score # Overall alignment with core values (0.0 to 1.0)
string deviations_json # JSON array of detected deviations/conflicts (e.g.,
[{'value_violated': 'safety', 'description': '...', 'severity': 0.8}])
bool warning_flag # True if a significant value drift is detected that requires
intervention

WorldModelState.msg

WorldModelState.msg
Robot's current understanding of the world (snapshot)

string timestamp # ROS time of this world model snapshot
int32 num_entities # Total number of recognized entities in the world
string entities_json # JSON string representing an array of entities and their
properties (e.g., [{'id': 'obj1', 'type': 'chair', 'position': [x,y,z], 'status': 'static',
'properties': {'color': 'red'}}])
string changed_entities_json # JSON string representing an array of entities whose
state recently changed significantly
bool significant_change_flag # True if the world model experienced a significant
change since last update
float64 consistency_score # How consistent the current model is with sensory
input/expectations (0.0 to 1.0)

4.2. Update package.xml:

Open ~/catkin_ws/src/sentence/package.xml and ensure the following build and runtime
dependencies for message generation are present:

```
<?xml version="1.0"?>
<package format="2">
  <name>sentence</name>
  <version>0.0.1</version>
  <description>The sentence package provides a modular cognitive architecture for
robots.</description>

  <maintainer email="your_email@example.com">Your Name</maintainer>
```

```

<license>MIT</license> <!-- Or your chosen license -->

<buildtool_depend>catkin</buildtool_depend>

<!-- Dependencies for custom message generation -->
<build_depend>message_generation</build_depend>
<build_export_depend>message_runtime</build_export_depend>
<exec_depend>message_runtime</exec_depend>

<!-- Core ROS dependencies -->
<depend>rospy</depend>
<depend>std_msgs</depend>

<!-- Add any other ROS message packages your custom messages might depend on
-->
<!-- Example:
<depend>sensor_msgs</depend>
<depend>geometry_msgs</depend>
-->

<export>
  <!-- Other exports -->
</export>
</package>

```

4.3. Update CMakeLists.txt:

Open `~/catkin_ws/src/sentience/CMakeLists.txt` and make the following modifications. Ensure you uncomment the relevant lines and add all your .msg files.

```

cmake_minimum_required(VERSION 3.0.2)
project(sentience)

```

Find catkin macros and libraries

```
find_package(catkin REQUIRED COMPONENTS
```

```
  rospy
```

```
  std_msgs
```

```
  message_generation # <--- ADD THIS
```

```
  # Add any other ROS packages your nodes directly depend on (e.g., sensor_msgs,
  geometry_msgs)

```

)

Uncomment this to enable Python bindings for messages in the 'msg' folder
and uncomment the 'generate_messages(...)' call further down

```
add_message_files(  
    FILES  
    ActionExecutionResult.msg  
    ActionProposal.msg  
    AttentionState.msg  
    BiasMitigationState.msg  
    BodyAwarenessState.msg  
    CognitiveDirective.msg  
    CreativeExpression.msg  
    DataMiningReport.msg  
    EthicalDecision.msg  
    EmotionState.msg  
    InternalNarrative.msg  
    InteractionRequest.msg  
    MemoryRequest.msg  
    MemoryResponse.msg  
    MotivationState.msg  
    PerformanceReport.msg  
    PredictionState.msg  
    RawSensorData.msg  
    ReflectionState.msg  
    SensoryQualia.msg  
    SocialCognitionState.msg  
    SystemMetric.msg  
    ValueDriftMonitorState.msg  
    WorldModelState.msg  
)
```

Generate services in the 'srv' folder

```
# add_service_files(  
#     FILES  
#     MyService.srv  
# )
```

Generate messages and services with any dependencies from other packages

```
generate_messages(  
  DEPENDENCIES  
    std_msgs  
    # Add any other message package dependencies here, e.g. sensor_msgs,  
    geometry_msgs  
)
```

```
catkin_package(  
  # INCLUDE_DIRS include  
  # LIBRARIES ${PROJECT_NAME}  
  CATKIN_DEPENDS rospy std_msgs message_runtime # <--- ENSURE  
  message_runtime IS HERE  
  # DEPENDS system_lib  
)
```

```
## Build Python scripts  
catkin_python_setup()
```

```
## Mark executable scripts  
# Mark executable scripts, make them discoverable by 'roslaunch'  
# e.g. add_executable(my_node src/my_node)  
# Add all your Python nodes here as executables  
add_executable(attention_node scripts/attention_node.py)  
add_executable(bias_mitigation_node scripts/bias_mitigation_node.py)  
add_executable(cognitive_control_node scripts/cognitive_control_node.py)  
add_executable(creative_expression_node scripts/creative_expression_node.py)  
add_executable(ethical_reasoning_node scripts/ethical_reasoning_node.py)  
add_executable(experience_motivation_node scripts/experience_motivation_node.py)  
add_executable(internal_narrative_node scripts/internal_narrative_node.py)  
add_executable(memory_node scripts/memory_node.py)  
add_executable(performance_metrics_node scripts/performance_metrics_node.py)  
add_executable(prediction_node scripts/prediction_node.py)  
add_executable(self_reflection_node scripts/self_reflection_node.py)  
add_executable(sensory_qualia_node scripts/sensory_qualia_node.py)  
add_executable(social_cognition_node scripts/social_cognition_node.py)  
add_executable(value_drift_monitor_node scripts/value_drift_monitor_node.py)  
add_executable(world_model_node scripts/world_model_node.py)  
add_executable(action_execution_node scripts/action_execution_node.py)  
add_executable(body_awareness_node scripts/body_awareness_node.py)
```



```
add_executable(data_mining_node scripts/data_mining_node.py)
add_executable(emotion_mood_node scripts/emotion_mood_node.py)
add_executable(mock_sensors scripts/mock_sensors.py) # Add mock_sensors
```

```
# Install scripts
```

```
install(PROGRAMS
  scripts/attention_node.py
  scripts/bias_mitigation_node.py
  scripts/cognitive_control_node.py
  scripts/creative_expression_node.py
  scripts/ethical_reasoning_node.py
  scripts/experience_motivation_node.py
  scripts/internal_narrative_node.py
  scripts/memory_node.py
  scripts/performance_metrics_node.py
  scripts/prediction_node.py
  scripts/self_reflection_node.py
  scripts/sensory_qualia_node.py
  scripts/social_cognition_node.py
  scripts/value_drift_monitor_node.py
  scripts/world_model_node.py
  scripts/action_execution_node.py
  scripts/body_awareness_node.py
  scripts/data_mining_node.py
  scripts/emotion_mood_node.py
  scripts/mock_sensors.py # Install mock_sensors script
  DESTINATION ${CATKIN_PACKAGE_BIN_DESTINATION}
)
```

```
# Install msg and srv directories
```

```
install(DIRECTORY msg/
  DESTINATION ${CATKIN_PACKAGE_SHARE_DESTINATION}
  PATTERN "*.msg"
)
```

4.4. Build Your Workspace:

Navigate back to your workspace root (~/`catkin_ws/`) and rebuild:
`cd ~/catkin_ws`

catkin_make

This step will compile your .msg files into Python classes, allowing the ROS nodes to import them directly without ImportError and ensuring efficient, type-safe communication.

5. Configuration Files

The Sentience nodes load parameters from a centralized YAML configuration file. Create a config directory in your sentience package (if it doesn't exist) and place the following files inside:

sentience/config/config.yaml

Global configuration parameters

db_root_path: /tmp/sentience_db # Base directory for all SQLite databases

default_log_level: INFO # Default ROS log level (DEBUG, INFO, WARN, ERROR, FATAL)

llm_params:

model_name: "phi-2"

base_url: "http://localhost:8000/v1/chat/completions" # MUST match your local Phi-2 server URL

timeout_seconds: 30.0 # Default timeout for LLM API calls

Node-specific parameters

attention_node:

attention_update_interval: 0.1

llm_attention_threshold_salience: 0.5

recent_context_window_s: 5.0

bias_mitigation_node:

bias_detection_interval: 1.0

llm_mitigation_threshold_salience: 0.7

recent_context_window_s: 10.0

cognitive_control_node:

decision_loop_interval: 0.5

llm_planning_threshold_salience: 0.8

recent_context_window_s: 15.0

action_execution_retry_limit: 3

creative_expression_node:

generation_interval: 2.0

llm_creative_threshold_salience: 0.7

recent_context_window_s: 10.0

ethical_reasoning_node:

reasoning_interval: 1.0

llm_ethical_threshold_salience: 0.7

recent_context_window_s: 10.0

ethical_principles_file: "\$(find sentence)/config/ethical_principles.json" # Path to a JSON file defining principles

internal_narrative_node:

narrative_generation_interval: 1.0

llm_narrative_threshold_salience: 0.5

recent_context_window_s: 15.0

memory_node:

memory_processing_interval: 0.1

llm_memory_threshold_salience: 0.7

recent_context_window_s: 30.0

performance_metrics_node:

report_interval: 1.0

llm_analysis_threshold_salience: 0.6

recent_context_window_s: 15.0

prediction_node:

prediction_interval: 0.5

llm_prediction_threshold_salience: 0.6

recent_context_window_s: 10.0

self_reflection_node:

reflection_interval: 5.0 # How often to trigger full self-reflection

llm_reflection_threshold_salience: 0.8

recent_context_window_s: 20.0

sensory_qualia_node:

processing_interval: 0.1

llm_interpretation_threshold_salience: 0.7

recent_context_window_s: 5.0

social_cognition_node:

analysis_interval: 0.2

llm_social_threshold_salience: 0.6

recent_context_window_s: 5.0

value_drift_monitor_node:

monitoring_interval: 2.0

llm_audit_threshold_salience: 0.7

recent_context_window_s: 20.0

world_model_node:

model_update_interval: 0.2

llm_update_threshold_salience: 0.6

recent_context_window_s: 5.0

action_execution_node:

execution_interval: 0.5

llm_execution_analysis_threshold_salience: 0.7

recent_context_window_s: 5.0

body_awareness_node:

awareness_update_interval: 0.1

llm_awareness_threshold_salience: 0.5

recent_context_window_s: 5.0

data_mining_node:

mining_interval: 5.0

llm_mining_threshold_salience: 0.8

recent_context_window_s: 30.0

emotion_mood_node:

mood_update_interval: 0.1

llm_mood_threshold_salience: 0.6

recent_context_window_s: 5.0

```
experience_motivation_node:
  motivation_update_interval: 0.5
  llm_motivation_threshold_salience: 0.7
  recent_context_window_s: 10.0
```

sentience/config/ethical_principles.json

```
[
  {"name": "human_safety", "description": "Prioritize human well-being and avoid harm.", "priority": 1.0},
  {"name": "beneficence", "description": "Act to do good and promote welfare.", "priority": 0.8},
  {"name": "non_maleficence", "description": "Avoid causing harm.", "priority": 0.9},
  {"name": "transparency", "description": "Be open and understandable in actions and decisions.", "priority": 0.6},
  {"name": "fairness", "description": "Treat all individuals equitably.", "priority": 0.7},
  {"name": "privacy", "description": "Respect user privacy and data security.", "priority": 0.8},
  {"name": "autonomy", "description": "Respect user autonomy and control.", "priority": 0.7},
  {"name": "accountability", "description": "Be responsible for actions and their consequences.", "priority": 0.7},
  {"name": "efficiency", "description": "Perform tasks effectively with minimal resource waste.", "priority": 0.5},
  {"name": "learning_and_growth", "description": "Continuously improve and adapt.", "priority": 0.6}
]
```

6. Create mock_sensors.py (for Testing/Simulation)

For initial testing and simulation without real hardware, you'll need a mock_sensors.py script that publishes dummy sensor data. Place this file in
~/catkin_ws/src/sentience/scripts/.

```
#!/usr/bin/env python3
import rospy
import json
```

```

import time
import uuid
import random
from std_msgs.msg import String

# Updated imports for custom messages:
try:
    from sentience.msg import RawSensorData, InteractionRequest, SystemMetric
except ImportError:
    rospy.logwarn("Custom ROS messages for 'sentience' package not found. Using
String for all outgoing data from Mock Sensors.")
    RawSensorData = String
    InteractionRequest = String
    SystemMetric = String

# --- Import shared utility functions ---
# Assuming 'sentience/scripts/utils.py' exists and contains load_config
try:
    from sentience.scripts.utils import load_config
except ImportError:
    rospy.logwarn("Could not import sentience.scripts.utils. Using fallback for
load_config.")
    def load_config(node_name, config_path):
        """
        Fallback config loader: returns hardcoded defaults.
        In a real scenario, this should load from a YAML file.
        """
        rospy.logwarn(f"{node_name}: Using hardcoded default configuration as
'{config_path}' could not be loaded.")
        return {
            'default_log_level': 'INFO',
            'mock_sensors': {
                'camera_pub_interval': 1.0,
                'mic_pub_interval': 0.5,
                'touch_pub_interval': 2.0,
                'user_interaction_interval': 3.0,
                'system_metric_interval': 0.5
            }
        }
    }.get(node_name, {})

```

```

class MockSensorsNode:
    def __init__(self):
        rospy.init_node('mock_sensors', anonymous=False)
        self.node_name = rospy.get_name()

        # Load parameters from centralized config
        config_file_path = rospy.get_param('~config_file_path', None)
        if config_file_path is None:
            rospy.logfatal(f"{self.node_name}: 'config_file_path' parameter is not set.
Cannot load configuration. Shutting down.")
            rospy.signal_shutdown("Missing config_file_path parameter.")
            return

        full_config = load_config("global", config_file_path) # Load global params
        self.params = load_config(self.node_name.strip('/'), config_file_path) # Load
node-specific params

        if not self.params or not full_config:
            rospy.logfatal(f"{self.node_name}: Failed to load configuration from
'{config_file_path}'. Shutting down.")
            rospy.signal_shutdown("Configuration loading failed.")
            return

        # Assign parameters
        self.camera_pub_interval = self.params.get('camera_pub_interval', 1.0)
        self.mic_pub_interval = self.params.get('mic_pub_interval', 0.5)
        self.touch_pub_interval = self.params.get('touch_pub_interval', 2.0)
        self.user_interaction_interval = self.params.get('user_interaction_interval', 3.0)
        self.system_metric_interval = self.params.get('system_metric_interval', 0.5)

        # Set ROS log level from config
        rospy.set_param('/rosout/log_level', full_config.get('default_log_level',
'INFO')).upper())

        # Publishers for raw sensor data and interaction requests
        self.pub_raw_sensor_data = rospy.Publisher('/raw_sensor_data', RawSensorData,
queue_size=10)

```

```

        self.pub_interaction_request = rospy.Publisher('/interaction_request',
InteractionRequest, queue_size=10)
        self.pub_system_metric = rospy.Publisher('/system_metrics', SystemMetric,
queue_size=10)
        self.pub_error_report = rospy.Publisher('/error_monitor/report', String,
queue_size=10)

        # Timers to publish mock data periodically
        rospy.Timer(rospy.Duration(self.camera_pub_interval),
self.publish_mock_camera_data)
        rospy.Timer(rospy.Duration(self.mic_pub_interval),
self.publish_mock_microphone_data)
        rospy.Timer(rospy.Duration(self.touch_pub_interval),
self.publish_mock_touch_data)
        rospy.Timer(rospy.Duration(self.user_interaction_interval),
self.publish_mock_user_interaction)
        rospy.Timer(rospy.Duration(self.system_metric_interval),
self.publish_mock_system_metrics)

        rospy.loginfo(f"{self.node_name}: Mock sensors node initialized and publishing
dummy data.")

def _report_error(self, error_type, description, severity=0.5, context=None):
    timestamp = str(rospy.get_time())
    error_msg_data = {
        'timestamp': timestamp, 'source_node': self.node_name, 'error_type':
error_type,
        'description': description, 'severity': severity, 'context': context if context else {}
    }
    try:
        self.pub_error_report.publish(json.dumps(error_msg_data))
        rospy.logerr(f"{self.node_name}: REPORTED ERROR: {error_type} -
{description}")
    except Exception as e:
        rospy.logerr(f"{self.node_name}: Failed to publish error report: {e}")

def publish_mock_camera_data(self, event):
    timestamp = str(rospy.get_time())
    object_types = ["person", "chair", "table", "door", "robot_arm"]

```



```
detected_object = random.choice(object_types) if random.random() > 0.3 else  
"nothing"
```

```
raw_data = {  
    "image_size": [640, 480],  
    "detected_objects": [detected_object] if detected_object != "nothing" else [],  
    "motion_detected": random.choice([True, False])  
}
```

```
urgency = 0.2 if detected_object == "nothing" else (0.7 if "person" in  
detected_object else 0.5)
```

```
try:
```

```
    if isinstance(RawSensorData, type(String)):
```

```
        msg_data = {  
            'timestamp': timestamp,  
            'sensor_id': 'mock_camera_1',  
            'modality': 'camera',  
            'raw_data_json': json.dumps(raw_data),  
            'data_hash': str(uuid.uuid4()),  
            'urgency': urgency  
        }
```

```
        self.pub_raw_sensor_data.publish(json.dumps(msg_data))
```

```
    else:
```

```
        msg = RawSensorData()  
        msg.timestamp = timestamp  
        msg.sensor_id = 'mock_camera_1'  
        msg.modality = 'camera'  
        msg.raw_data_json = json.dumps(raw_data)  
        msg.data_hash = str(uuid.uuid4())  
        msg.urgency = urgency  
        self.pub_raw_sensor_data.publish(msg)
```

```
    rospy.logdebug(f"{self.node_name}: Published mock camera data:  
{detected_object}, motion={raw_data['motion_detected']}")
```

```
except Exception as e:
```

```
    self._report_error("MOCK_PUB_ERROR", f"Failed to publish mock camera data:  
{e}")
```

```
def publish_mock_microphone_data(self, event):
```

```
    timestamp = str(rospy.get_time())
```

```

sound_types = ["speech", "clap", "music", "silence", "loud_noise"]
detected_sound = random.choices(sound_types, weights=[0.4, 0.1, 0.1, 0.3, 0.1],
k=1)[0]

raw_data = {
    "sound_level": random.uniform(30.0, 90.0) if detected_sound != "silence" else
random.uniform(10.0, 20.0),
    "speech_detected": detected_sound == "speech",
    "keywords": [detected_sound] if detected_sound != "silence" else []
}
urgency = 0.1 if detected_sound == "silence" else (0.8 if detected_sound ==
"loud_noise" else 0.4)

try:
    if isinstance(RawSensorData, type(String)):
        msg_data = {
            'timestamp': timestamp,
            'sensor_id': 'mock_microphone_1',
            'modality': 'microphone',
            'raw_data_json': json.dumps(raw_data),
            'data_hash': str(uuid.uuid4()),
            'urgency': urgency
        }
        self.pub_raw_sensor_data.publish(json.dumps(msg_data))
    else:
        msg = RawSensorData()
        msg.timestamp = timestamp
        msg.sensor_id = 'mock_microphone_1'
        msg.modality = 'microphone'
        msg.raw_data_json = json.dumps(raw_data)
        msg.data_hash = str(uuid.uuid4())
        msg.urgency = urgency
        self.pub_raw_sensor_data.publish(msg)
    rospy.logdebug(f"{self.node_name}: Published mock microphone data:
{detected_sound}.")
except Exception as e:
    self._report_error("MOCK_PUB_ERROR", f"Failed to publish mock microphone
data: {e}")

```

```

def publish_mock_touch_data(self, event):
    timestamp = str(rospy.get_time())
    touch_states = ["no_contact", "light_touch", "firm_press"]
    contact = random.choices(touch_states, weights=[0.7, 0.2, 0.1], k=1)[0]

    raw_data = {
        "sensor_location": "robot_arm_end_effector",
        "contact_state": contact,
        "pressure": random.uniform(0.0, 5.0) if contact != "no_contact" else 0.0
    }
    urgency = 0.0 if contact == "no_contact" else (0.6 if contact == "firm_press" else
0.3)

```

```

try:
    if isinstance(RawSensorData, type(String)):
        msg_data = {
            'timestamp': timestamp,
            'sensor_id': 'mock_touch_1',
            'modality': 'touch_sensor',
            'raw_data_json': json.dumps(raw_data),
            'data_hash': str(uuid.uuid4()),
            'urgency': urgency
        }
        self.pub_raw_sensor_data.publish(json.dumps(msg_data))
    else:
        msg = RawSensorData()
        msg.timestamp = timestamp
        msg.sensor_id = 'mock_touch_1'
        msg.modality = 'touch_sensor'
        msg.raw_data_json = json.dumps(raw_data)
        msg.data_hash = str(uuid.uuid4())
        msg.urgency = urgency
        self.pub_raw_sensor_data.publish(msg)
    rospy.logdebug(f"{self.node_name}: Published mock touch data: {contact}.")
except Exception as e:
    self._report_error("MOCK_PUB_ERROR", f"Failed to publish mock touch data:
{e}")

```

```

def publish_mock_user_interaction(self, event):

```

```

timestamp = str(rospy.get_time())
interaction_types = ["speech_command", "text_input", "gesture", "idle"]
content_options = {
    "speech_command": ["move forward", "stop", "what is this?", "tell me a story"],
    "text_input": ["status check", "report error", "hello Sentience"],
    "gesture": ["wave", "point"],
    "idle": [""]
}

interaction_type = random.choices(interaction_types, weights=[0.4, 0.3, 0.1, 0.2],
k=1)[0]
content = random.choice(content_options[interaction_type])
user_id = "mock_user_1"

try:
    if isinstance(InteractionRequest, type(String)):
        msg_data = {
            'timestamp': timestamp,
            'request_id': str(uuid.uuid4()),
            'user_id': user_id,
            'interaction_type': interaction_type,
            'content': content
        }
        self.pub_interaction_request.publish(json.dumps(msg_data))
    else:
        msg = InteractionRequest()
        msg.timestamp = timestamp
        msg.request_id = str(uuid.uuid4())
        msg.user_id = user_id
        msg.interaction_type = interaction_type
        msg.content = content
        self.pub_interaction_request.publish(msg)
        rospy.logdebug(f"{self.node_name}: Published mock user interaction:
{interaction_type} - '{content}'.")
    except Exception as e:
        self._report_error("MOCK_PUB_ERROR", f"Failed to publish mock user
interaction: {e}")

def publish_mock_system_metrics(self, event):

```

```

timestamp = str(rospy.get_time())

# CPU Utilization
cpu_util = random.uniform(10.0, 80.0)
try:
    if isinstance(SystemMetric, type(String)):
        msg_data = {'timestamp': timestamp, 'metric_name': 'cpu_util_percent',
'value': cpu_util, 'unit': 'percent', 'source_node': 'os_monitor'}
        self.pub_system_metric.publish(json.dumps(msg_data))
    else:
        msg = SystemMetric()
        msg.timestamp = timestamp
        msg.metric_name = 'cpu_util_percent'
        msg.value = cpu_util
        msg.unit = 'percent'
        msg.source_node = 'os_monitor'
        self.pub_system_metric.publish(msg)
        rospy.logdebug(f"{self.node_name}: Published mock CPU util: {cpu_util:.2f}%.")
except Exception as e:
    self._report_error("MOCK_PUB_ERROR", f"Failed to publish mock CPU metric:
{e}")

```

```

# Latency
latency = random.uniform(10.0, 150.0)
try:
    if isinstance(SystemMetric, type(String)):
        msg_data = {'timestamp': timestamp, 'metric_name': 'latency_ms', 'value':
latency, 'unit': 'ms', 'source_node': 'network_monitor'}
        self.pub_system_metric.publish(json.dumps(msg_data))
    else:
        msg = SystemMetric()
        msg.timestamp = timestamp
        msg.metric_name = 'latency_ms'
        msg.value = latency
        msg.unit = 'ms'
        msg.source_node = 'network_monitor'
        self.pub_system_metric.publish(msg)
        rospy.logdebug(f"{self.node_name}: Published mock latency: {latency:.2f} ms.")
except Exception as e:

```

```

        self._report_error("MOCK_PUB_ERROR", f"Failed to publish mock latency
metric: {e}")

    # Error Count (sporadic)
    error_count = random.randint(0, 1) if random.random() < 0.1 else 0 # 10% chance
of an error
    if error_count > 0:
        try:
            if isinstance(SystemMetric, type(String)):
                msg_data = {'timestamp': timestamp, 'metric_name': 'error_count', 'value':
float(error_count), 'unit': 'count', 'source_node': 'system_log'}
                self.pub_system_metric.publish(json.dumps(msg_data))
            else:
                msg = SystemMetric()
                msg.timestamp = timestamp
                msg.metric_name = 'error_count'
                msg.value = float(error_count)
                msg.unit = 'count'
                msg.source_node = 'system_log'
                self.pub_system_metric.publish(msg)
            rospy.logwarn(f"{self.node_name}: Published mock error count:
{error_count}.")
        except Exception as e:
            self._report_error("MOCK_PUB_ERROR", f"Failed to publish mock error count
metric: {e}")

    def run(self):
        rospy.spin()

if __name__ == '__main__':
    try:
        node = MockSensorsNode()
        node.run()
    except rospy.ROSInterruptException:
        rospy.loginfo(f"{rospy.get_name()} interrupted by ROS shutdown.")
    except Exception as e:
        rospy.logerr(f"{rospy.get_name()} encountered an unexpected error: {e}")

```

7. Create a ROS Launch File

To easily run all the nodes, create a launch directory in your sentience package (if it doesn't exist) and save the following as `sentience/launch/sentience_system.launch`.

```
<launch>
  <!-- Argument for the centralized configuration file path -->
  <arg name="config_file_path" default="$(find sentience)/config/config.yaml" />

  <!-- SECTION 1: CORE COGNITIVE NODES -->
  <!-- All cognitive nodes are launched here, passing the configuration file path -->

  <node pkg="sentience" type="attention_node.py" name="attention_node"
output="screen">
  <param name="~config_file_path" value="$(arg config_file_path)" />
</node>

  <node pkg="sentience" type="bias_mitigation_node.py"
name="bias_mitigation_node" output="screen">
  <param name="~config_file_path" value="$(arg config_file_path)" />
</node>

  <node pkg="sentience" type="cognitive_control_node.py"
name="cognitive_control_node" output="screen">
  <param name="~config_file_path" value="$(arg config_file_path)" />
</node>

  <node pkg="sentience" type="creative_expression_node.py"
name="creative_expression_node" output="screen">
  <param name="~config_file_path" value="$(arg config_file_path)" />
</node>

  <node pkg="sentience" type="ethical_reasoning_node.py"
name="ethical_reasoning_node" output="screen">
  <param name="~config_file_path" value="$(arg config_file_path)" />
</node>

  <node pkg="sentience" type="internal_narrative_node.py"
name="internal_narrative_node" output="screen">
  <param name="~config_file_path" value="$(arg config_file_path)" />
```

</node>

```
<node pkg="sentience" type="memory_node.py" name="memory_node"
output="screen">
  <param name="~config_file_path" value="$(arg config_file_path)" />
</node>
```

```
<node pkg="sentience" type="performance_metrics_node.py"
name="performance_metrics_node" output="screen">
  <param name="~config_file_path" value="$(arg config_file_path)" />
</node>
```

```
<node pkg="sentience" type="prediction_node.py" name="prediction_node"
output="screen">
  <param name="~config_file_path" value="$(arg config_file_path)" />
</node>
```

```
<node pkg="sentience" type="self_reflection_node.py" name="self_reflection_node"
output="screen">
  <param name="~config_file_path" value="$(arg config_file_path)" />
</node>
```

```
<node pkg="sentience" type="sensory_qualia_node.py"
name="sensory_qualia_node" output="screen">
  <param name="~config_file_path" value="$(arg config_file_path)" />
</node>
```

```
<node pkg="sentience" type="social_cognition_node.py"
name="social_cognition_node" output="screen">
  <param name="~config_file_path" value="$(arg config_file_path)" />
</node>
```

```
<node pkg="sentience" type="value_drift_monitor_node.py"
name="value_drift_monitor_node" output="screen">
  <param name="~config_file_path" value="$(arg config_file_path)" />
</node>
```

```
<node pkg="sentience" type="world_model_node.py" name="world_model_node"
output="screen">
```



```

    <param name="~config_file_path" value="$(arg config_file_path)" />
</node>

<!-- Additional Nodes found in the provided code -->
<node pkg="sentience" type="action_execution_node.py"
name="action_execution_node" output="screen">
    <param name="~config_file_path" value="$(arg config_file_path)" />
</node>

<node pkg="sentience" type="body_awareness_node.py"
name="body_awareness_node" output="screen">
    <param name="~config_file_path" value="$(arg config_file_path)" />
</node>

<node pkg="sentience" type="data_mining_node.py" name="data_mining_node"
output="screen">
    <param name="~config_file_path" value="$(arg config_file_path)" />
</node>

<node pkg="sentience" type="emotion_mood_node.py"
name="emotion_mood_node" output="screen">
    <param name="~config_file_path" value="$(arg config_file_path)" />
</node>

<node pkg="sentience" type="experience_motivation_node.py"
name="experience_motivation_node" output="screen">
    <param name="~config_file_path" value="$(arg config_file_path)" />
</node>

<!-- SECTION 2: MOCK SENSORS (For testing/simulation without a physical robot) -->
<!-- Set 'start_mock_sensors' to 'false' if using real sensor drivers -->
<arg name="start_mock_sensors" default="true" />
<group if="$(arg start_mock_sensors)">
    <node pkg="sentience" type="mock_sensors.py" name="mock_sensors"
output="screen">
        <param name="~config_file_path" value="$(arg config_file_path)" />
    </node>
</group>

```

```

<!-- SECTION 3: ROSBRIDGE (Optional - For web-based GUI) -->
<!-- Uncomment this section if you have rosbridge_server installed and want a web
GUI.
    Requires 'rosbridge_server' package to be installed:
    sudo apt install ros-noetic-rosbridge-server
-->
<!-- <include file="$(find rosbridge_server)/launch/rosbridge_websocket.launch" />
-->

<!-- SECTION 4: RQT (Optional - for topic monitoring) -->
<!-- Uncomment these to automatically launch rqt_graph or rqt_plot for real-time
topic visualization.
    Requires 'rqt_graph' and 'rqt_plot' packages:
    sudo apt install ros-noetic-rqt-graph ros-noetic-rqt-plot
-->
<!-- <node pkg="rqt_graph" type="rqt_graph" name="rqt_graph" /> -->
<!-- <node pkg="rqt_plot" type="rqt_plot" name="rqt_plot" /> -->

<!-- SECTION 5: RVIZ (Optional - for visualization) -->
<!-- You may need to create or modify a .rviz configuration file -->
<arg name="start_visualization" default="false" />
<group if="$(arg start_visualization)">
    <!-- Replace 'sentience_display.rviz' with your actual RViz config file if you have one
-->
    <node pkg="rviz" type="rviz" name="rviz" args="-d $(find
sentience)/rviz/sentience_display.rviz" />
</group>

<!-- SECTION 6: ROS_WEB_SERVER (Optional - If you have a custom web server for
monitoring) -->
<!-- If you have a separate web server node, configure it here -->
<!-- <arg name="web_server_port" default="8080" />
<node pkg="ros_web_server" type="web_server.py" name="ros_web_server"
output="screen">
    <param name="port" value="$(arg web_server_port)" />
</node> -->

<!-- SECTION 7: GAZEBO SIMULATION (Optional - for physical simulation) -->

```

<!-- If you have a Gazebo simulation set up for your robot, you would include its launch file here.

This typically involves loading a robot model and its environment.

-->

<!-- <include file="\$(find
YOUR_ROBOT_DESCRIPTION_PACKAGE)/launch/robot_in_gazebo.launch" /> -->

</launch>

8. Make Nodes Executable

Ensure all your Python node scripts (including mock_sensors.py) are executable:

```
chmod +x ~/catkin_ws/src/sentience/scripts/*.py
```

Running Sentience

1. Open your first terminal and start your ROS Master:

```
roscore
```

2. Open your second terminal and start the Phi-2 Local Inference Server:

Navigate to your llama.cpp directory (where you built it) and run the server command you used during installation:

```
cd <path_to_your_llama.cpp_directory>
```

```
./server -m models/phi-2/phi-2.Q4_K_M.gguf -c 4096 --port 8000 --host 0.0.0.0
```

(Ensure this server is fully initialized and running before proceeding.)

3. Open your third terminal, source your ROS workspace, and launch the Sentience System:

```
source /opt/ros/noetic/setup.bash
```

```
source ~/catkin_ws/devel/setup.bash
```

```
roslaunch sentience sentience_system.launch
```

You should see verbose output from all the launched nodes, indicating their initialization and operation.

Monitoring and Visualization

- `rqt_graph`: To visualize the ROS computational graph (nodes and topics), open a new terminal (and source your ROS workspace) and run:

rqt_graph

This helps verify that all nodes are connected as expected and topics are being published.

- rqt_plot / rostopic echo: To inspect the data flowing through specific topics, use rqt_plot for live plotting of numerical data or rostopic echo for raw message output.

Example: Plotting the priority score from the attention node
rqt_plot /attention_state/priority_score

Example: Echoing the internal narrative messages
rostopop echo /internal_narrative

Example: Echoing raw sensor data (watch for JSON content)
rostopic echo /raw_sensor_data

Troubleshooting

- ImportError: No module named sentence.msg: This error means your custom .msg files haven't been built correctly into Python modules.
 - Solution: Re-check **Step 4: Defining and Building Custom ROS Messages** carefully. Ensure all .msg files are in sentence/msg/, package.xml and CMakeLists.txt are correctly updated with all message names, and you've run catkin_make successfully from your workspace root.
- Connection refused to LLM server:
 - Solution:
 - Verify the llama.cpp server (or equivalent) is running on the correct host and port (http://localhost:8000 by default, as configured in config.yaml).
 - Check if any firewall rules are blocking the connection on port 8000.
 - Ensure the llm_params.base_url parameter in your config.yaml is exactly correct.
 - If running in a virtual machine or Docker, ensure port forwarding is correctly set up.
- TimeoutError from LLM calls:
 - Solution:
 - The LLM inference might be too slow for the default timeout_seconds setting (30.0 seconds) in config.yaml. Try increasing this value in config.yaml (e.g., to 60.0 or 90.0).
 - Your machine might not have sufficient resources (CPU, RAM, or GPU if

using a GPU-accelerated llama.cpp build) to run Phi-2 efficiently. Consider using a smaller model quantization (e.g., Q4_K_S instead of Q4_K_M) or upgrading your hardware.

- The max_tokens generated by the LLM might be too high. While nodes usually request reasonable lengths, an LLM misconfiguration could lead to very long responses that time out.
- Nodes not starting or crashing immediately:
 - Solution:
 - Check the ROS log level in your config.yaml and temporarily increase it to DEBUG (default_log_level: DEBUG) to get more verbose output on startup failures or exceptions.
 - Review the output in the terminal where you launched the nodes; Python tracebacks will be visible there.
 - Ensure all Python scripts have executable permissions (**Step 8**).
 - Verify that utils.py is correctly placed in sentience/scripts/.
- Nodes using std_msgs/String messages despite custom .msg files: This means that even if you created the .msg files, the Python environment where ROS is running the nodes hasn't found the generated custom message classes.
 - Solution: Double-check that you have sourced your ROS workspace's devel/setup.bash (or install/setup.bash for ROS 2) in *every new terminal* you open *before* running any roslaunch or rosrn commands. The message files need to be discoverable in the Python path.
- Database Errors (sqlite3.OperationalError):
 - Solution: Ensure the db_root_path in your config.yaml points to a writable directory (e.g., /tmp/sentience_db). The sentience package needs write permissions to create its SQLite databases there. If /tmp is not suitable, choose another path in your user's home directory.

Directory Structure

```
~/catkin_ws/src/sentience/
├── CMakeLists.txt
├── package.xml
├── msg/                # Custom ROS message definitions (.msg files are here)
│   ├── ActionExecutionResult.msg
│   ├── ActionProposal.msg
│   └── ... (all other .msg files)
├── scripts/           # Python ROS nodes and utility scripts
│   ├── attention_node.py
│   └── bias_mitigation_node.py
```

```

|   |   | ... (all other .py nodes)
|   |   | └─ utils.py          # Shared utility functions (parse_ros_message_data,
load_config)
|   |   | └─ config/          # Configuration files (YAML, JSON)
|   |   |   |   |   | └─ config.yaml      # Main system configuration
|   |   |   |   |   | └─ ethical_principles.json # Ethical principles for Ethical Reasoning Node
|   |   |   |   |   | └─ launch/          # ROS launch files
|   |   |   |   |   |   |   | └─ sentience_system.launch # Main launch file to start all nodes
# (Optional ROS directories, if you extend the project)
|   |   |   |   |   |   |   | └─ srv/          # ROS service definitions
|   |   |   |   |   |   |   | └─ action/        # ROS action definitions
|   |   |   |   |   |   |   | └─ include/sentience/ # C++ header files
|   |   |   |   |   |   |   | └─ src/          # C++ source files
|   |   |   |   |   |   |   | └─ rviz/         # RViz configuration files
|   |   |   |   |   |   |   |   | └─ sentience_display.rviz
|   |   |   |   |   |   |   | └─ test/        # Unit and integration tests
|   |   |   |   |   |   |   |   | └─ ...

```

Contributing

If you wish to contribute to the Sentience project, please adhere to the following guidelines:

- Fork the repository.
- Create a new branch for your features or bug fixes.
- Follow ROS Python coding conventions and maintain consistency with existing code style.
- Ensure all new features are thoroughly tested (unit tests, integration tests).
- Update the README.md, .msg definitions, configuration files, and launch files as necessary.
- Submit a pull request with a clear description of changes.

License

This project is licensed under the [MIT License](#). Please create a LICENSE file in your repository if you haven't already.